

The Strange Experience Argument for Modest Conceptual Functionalism

Jackson Kernion · UC Berkeley

1 · The Project: Are Phenomenal Concepts Functional Concepts?

We can think of the mind in *phenomenal* terms, in terms of what it feels like to have mental states.

We can think of the mind in *functional* terms, in terms of its causal organization and causal architecture.

<i>Phenomenal</i>		<i>Functional</i>
P ₁	← ? →	F ₁
P ₂	← ? →	F ₂
P ₃	← ? →	F ₃
⋮		⋮
P _n	← ? →	F _n

- Different ways of understanding the functional
- Mean to pick these out ostensively.
- Big Picture: how-it-feels vs. how-it-works.
- Example: your color experience, your visual cortex

Central question: How are these two ways of thinking related? Answered by mapping out a priori entailments...

- “Reading off”
- Do phenomenal facts have a functional upshot?
- How to do this? Inspect Individual entailments: example cases, conceivability tests
- Ex: We can ask both whether P1 entails F1.
- To test: can we conceive of P1 without F1? If not, P1 entail F1.
- Two directions of entailment.
- Different views: different ways of characterizing map of entailments.

STANDARD OPTIONS

Ambitious Conceptual Functionalism: for every phenomenal description, there's a matching functional description that both entails it, and is entailed by it.

- *Big Problem #1 (Absent qualia):* Functional descriptions can't, a priori, settle *whether* some system has phenomenal states.
- *Big Problem #2 (Inverted qualia):* Assuming some system *is* phenomenally conscious, functional descriptions can't, a priori, settle the precise *character* of its phenomenal states.

Conceptual Dualism: The phenomenal is conceptually independent of the functional. Phenomenal descriptions never entail functional descriptions, and vice versa.

"Phenomenal concepts are conceptually irreducible in this sense: they neither a priori imply, nor are implied by, physical-functional concepts."

"Phenomenal concepts are conceptually independent of physical-functional descriptions..."
—Brian Loar, "Phenomenal States" (1990)

SHARED WIDELY: metaphysical dualists, empirical functionalism, anti-realists

Modest Conceptual Functionalism: (At least some) phenomenal concepts have built-in functional constraints. So while functional descriptions don't entail phenomenal descriptions, (at least some) phenomenal descriptions entail functional descriptions.

BASIC MOTIVATION: We apply phenomenal properties to physical systems

NOT LIKE: refrigerators don't have uncles.

UPSHOT: skeptical of this talk of subjective concepts.

Where to look for phenomenal-to-functional entailments? I see two promising sources:

- *Subject-based constraints:* The idea of a phenomenal subject might contain within it an idea of how phenomenal subjects have to be structured, in functional terms. (It's telling that we first need to identify *candidate* experience-havers before we then go about *subtracting* their phenomenology in absent qualia cases.)
- *Phenomenal-kind-based constraints:* Different phenomenal kinds (e.g. color experience, auditory experience) might, necessarily, meet kind-specific functional constraints. (It's telling that we think red experience/green experience phenomenal swaps are, intuitively, more intelligible than, say, red experience/anxiety phenomenal swaps.)
- *Reveals: the structure of phenomenal variation for each phenomenal kind has to be matched by the structure of functional variation in any underlying state.* (Dissonant qualia cases can, I think, be used to uncover such constraints.)

How is this non-obvious? Attending carefully to a set of issues that have received little attention, clarifying the structure of the problem, demonstrating potentially under-appreciated consequences.

WHY DOES THIS ISSUE MATTER?

- *The intelligibility of physicalism*: Even if physicalism isn't true, it at least seems like an intelligible view—more intelligible than the view that phenomenal consciousness is the rate of economic growth in Bulgaria. Modest Conceptual Functionalism can make sense of this.
- *The tractability of a science of consciousness*: If Conceptual Dualism were true, it would make a science of consciousness impossible. Confronted with two different hypotheses about the actual correlation between phenomenal properties and functional properties, we wouldn't know how to bring 'objective', 'third-personal' evidence to bear on the issue. Phenomenal-to-functional entailments can provide 'guardrails' for our empirical project.

The rest of the talk: pursue the subject-based path, looking at particular phenomenal concept/functional concept pair.

CHECK TIME: 4:20

2 · P-Consciousness & A-Consciousness

TWO CONCEPTS OF CONSCIOUSNESS

Ned Block, famously, distinguished between P-consciousness and A-Consciousness:

A mental state is *phenomenally (P-) conscious* just in case it feels like something to be in that state.

- Hard to define non-circular way

A mental state is *access (A-) conscious* just in case it is directly available for arbitrary use within our cognitive economy, it's uniquely well-positioned to directly influence one's thoughts and behaviors.

- Working memory, RAM.
- Purpose is to keep information 'live', readily available for reference or further processing.
- PRECISE CHARACTERIZATION IS DISPUTED
- UNNECESSARILY RESTRICTIVE: by focusing on specific information-consuming capacities (e.g. speech production, abstract reasoning).
- EXAMPLE: Organisms that lack language, or rational capacities. Still surface information.
- ALSO: a state can be *available* for reasoning or verbal report, without *actually* being used for those purposes.
- PICKS OUT the feature of cognitive systems that *underlies* or *explains* how a mental state can be used in reasoning or in the deliberate production of behavior
- CAN ALSO BE TOO PERMISSIVE: by failing to distinguish *directly* vs. *indirectly* available.
- EX: beliefs are available to central cognition in and indirect sense. Need to be triggered or deliberately called up for them to 'come online', to be available in a more direct sense.

- This distinction between directly available ‘online’ states and indirectly available ‘offline’ states can be spelled out in different ways, but, for now, just know that some such distinction needs to be made for there to be a meaningful distinction between a-conscious states and non-a-conscious states
- GRAMMATICAL POINTS
- First, phenomenal consciousness and access consciousness are categories we apply to mental states in virtue of their being conscious *at all*, without regard to their determinate character or specific content.
- Second, we typically take phenomenal consciousness and access consciousness to be *relational* properties, relating individual mental states to the subject to whom they belong.

DISTINCT CONCEPTS?

- Lots of philosophers think there’s an *actual* overlap of p-consc and a-consc. But is it necessary?

A-to-P entailment: Is A-consciousness without P-consciousness conceivable? (Yes.)

P-to-A entailment: Is P-consciousness without A-consciousness conceivable?

“Block’s distinction between access consciousness and phenomenal consciousness (or experience) is very useful. There is clearly a conceptual distinction here, as illustrated by the facts that: (1) **one can imagine access without experience and vice versa**; (2) access can be observed straightforwardly, whereas experience cannot; and, most important, (3) access consciousness seems clearly amenable to cognitive explanation, whereas phenomenal consciousness is quite perplexing in this regard.” (Emphasis added)

—David Chalmers, “Availability: The cognitive basis of experience?” (1997)

3 · Strange Experiences, Strange Subjects, and More

Strange Experiences: mental states that are P-conscious for some subject, yet functionally isolated from that subject’s A-consciousness.

EXAMPLE: STRANGE PAIN

You’re experiencing a sharp, piercing pain in your lower back. But there’s something strange about this pain. It doesn’t feature in your mental life in the way that ‘normal’ experiences do. It’s causally isolated from your other ‘conscious’ mental processes:

- You’re unable to determine the pain’s location (even though you *feel* it in lower back)
- You can’t describe its qualities (even though it distinctly *feels* sharp and piercing)
- You’re unable to notice that your back feels unusual (even though it *feels* very unusual)

You could go about your day as you normally would without this pain ‘making a difference’ or distracting you. You could hear this very description and think, “Wow, what a horrible scenario!”, without realizing that this describes precisely your own situation.

It’s not the absence of pain *behavior* (or any outwardly observable effects) that makes this pain strange. Instead, what’s strange is that the pain doesn’t show up to *you*: it has no ‘inwardly’ observable effects. You’re not merely doing a great job of masking your pain from others or from yourself; rather, *your* pain is, in a sense, masked from *you*.

Pain may be particularly evocative, but the strangeness, here, doesn't hinge on the experience being a *pain* rather than some other kind of experience with a different phenomenal character. I could just as easily describe a *strange tickle* (i.e. an uncomfortable tickle that you *feel* on your left abdomen but is functionally isolated) or a *strange smell* (i.e. a pungent rotting corpse odor that you *experience* but is functionally isolated).

PROGRESSIVELY STRANGER CASES

- *Strange Experience Collection*: A subject has more than one strange experience at once (e.g. strange pain + strange tickle + strange smell).
- *Strange Modality*: An entire sensory modality of a subject goes strange (e.g. All visual experience is made strange).
- *Strange Subjects*: Every P-conscious state of a subject goes strange (i.e. The opposite of a philosophical zombie: all P-conscious states and no A-conscious states).

STRANGE SUBJECTS

Strange subjects are the opposite of (philosophical) zombies. And just as you can try to conceive of your zombie counterpart, you can also try to conceive of your strange subject counterpart (or, strange twin). Whereas my zombie twin is a physical and functional duplicate that lacks p-consciousness, my strange twin will be a phenomenal duplicate that lacks a-consciousness.

Remember, your strange twin will, by definition, share *all* of your p-conscious states. So your total p-conscious experience should be subjectively indistinguishable from your strange twin's. Presuming that you know exactly *what it is like* to be you, you should already know exactly *what it is like* to be your strange twin. It should therefore be fairly straightforward to conceive of *this* aspect of strange subjects.

Difficulty emerges when you try to add in a global constraint on the a-consciousness of strange subjects. Strange subjects will share their phenomenology with us, but we’re stipulating that none of their p-conscious states are a-conscious. So to conceive of your strange twin, just try to subtract a-consciousness from your own case while leaving all p-conscious aspect of your mental states intact.

SKEPTICISM ABOUT CONCEIVABILITY OF STRANGE SUBJECTS

1. **The Subjective Unity Worry**: Is there really any subject left remaining if the various experiences that ‘come together’ in a subject can’t bear any functional relations to one another and to the subject which owns them?

If a strange subject’s every experience is isolated from itself

DIFFICULT to project ourselves into the strange subject's point of view.

Any point of view left remaining if nothing is 'presented' to that point of view?

Is there really any subject left remaining if the various experiences that 'come together' in a subject can't bear any functional relations to one another and to the subject who owns them? To have a subject at all, various experiences may have to be functionally integrated in order to form a single cohesive subject.

2. **The Mental Action Worry:** We typically think of (even purely phenomenal) minds as capable of mental action: they make decisions, form judgments, initiate action, and rehearse mental routines. Would any of this be possible for a subject that lacked any and all functional organization?

Even when thinking of a purely phenomenal mind (e.g. something like a Cartesian Ego), we think about such minds as capable of mental action.

Minds can *do* things: they make DECISION; they form JUDGMENTS; they REHEARSE mental routines—ALL this can be understood entirely from the subject's own internal perspective. It's not clear that *any* of this would be possible for a subject that lacked any and all functional organization.

EX: perform a calculation or pursue some action, mental or otherwise?

EX: *feel* like there's some mental activity going on without there actually being any corresponding mental activity?

3. **The Temporal Relations Worry:** We typically think that the experiences of a subject at one time are causally related to the experiences of the same subject at time a later time. Could a strange subject be temporally structured in this way? If not, can we make sense of a single strange subject having a sequence of experiences?

When a strange subject feels a painful sensation (e.g. the sensation of a punch to the gut), then has a visual experience as of the cause of that sensation (e.g. seeing (or seeming to see) their friend Frank punching them), then has an emotional 'response' (e.g. anger at Frank), then feels a desire to take a particular action (e.g. a desire to punch Frank back), and then experiences the performance of that action (e.g. the punching of Frank)—could any of these experiences be causally related (given the constraint on functional organization)?

If not, can we make sense of a single strange subject having a sequence of experiences like this? Or must strange subjects, as a result of lacking the functional structure which would maintain them through time, exist only instantaneously?

NOTE: These 'global' Strange Subject worries should have 'local' analogues in other cases.

Structured Subject Intuition: Conceiving of a phenomenal mind already involves the conception of some kind of underlying functional structure in which various experiences are embedded and through which they bear relations to one another. We don't conceive of the ebb and flow of experience as being structurally/functionally unconstrained—we necessarily apply a structural/functional lens when thinking about the collection of experiences which comprise a given subject.

ADDITIONAL STRANGE CASES

- *Strange Modality Expansion:* A subject has an 'extra' Strange Modality. Consider, for instance, having a Strange Modality which gives you inaccessible bat sonar experiences.
- *Strange Swapping:* Two different subjects are A-conscious of each *other's* P-conscious states (or vice versa). For instance, consider a case in which 'you' are currently experiencing the life of someone else while also thinking about and making decisions about your own life.
- *Strange God:* A subject who 'has' every experience in the world (i.e. your experiences, every bat's experiences, etc.) but is unable to access any of them. Similarly, consider a subject who has every experiences but can only access the experiences of a single organism (e.g. you experience everything, but are only A-conscious of 'your' P-conscious states.)

CHECK TIME: 4:45

4 • Can Strange Experiences Be Made Palatable?

If strange experiences are not conceivable, then Conceptual Dualism will lose the justification necessary for their denial of P-to-A a priori entailment. Consequently, it's unclear if there really is a clean distinction between thinking of 'consciousness' in a how-it-feels vs. how-it-works way.

'RADICAL' A PRIORI SKEPTICISM

Objection: "Why be so sure that there's some fact of the matter about whether strange subjects are conceivable? It's a mistake to insist on a clear verdict in such problem cases."

Reply: This approach can't be used in support of Conceptual Dualism. It proves too much: it undermines the notions of conceivability/the a priori that Conceptual Dualists rely on.

DIFFERENT NOTIONS OF ACCESS

Objection: "We can distinguish between mental states that are *actually accessed* from those which are *merely accessible*. And if we understand 'A-conscious' to mean 'actually accessed', then it's easier to accept that there can be P-conscious states that don't become A-conscious."

Reply: The basic problem with this kind of approach is that it merely relocates, rather than removes, the challenge. Replacing an '*all P-conscious states must be accessed*' constraint with an '*all P-conscious states must be accessible*' constraint does nothing to remove the motivating worry.

DIMMED EXPERIENCE

Objection: “There are *actual* cases of P-consciousness without A-consciousness: cases of experiences at the periphery of our phenomenology. We typically think such ‘dimmed’ experiences are still *experiences*, but they may not be A-conscious in the usual way.”

Reply: Dimmed experiences look like ‘unaccessed but still accessible’ mental states. And it would be much less clear that these were really experiences if they were totally inaccessible.

UNIMAGINABLE, NOT INCONCEIVABLE

Objection: “Are we sure that we run into any tension just in *conceiving* of inaccessible experiences? It may be an act of *imagination* that causes problems: we try to simulate what a strange experience would feel like, and then when we ‘check on’ this simulated experience, we, unsurprisingly, always find such experiences to be A-conscious. Such ‘checking’ requires A-consciousness”

Reply: Our concept of experience can be applied without actively imagining or conceiving of any determinate kind of experience. And we get the same conceivability verdict in cases where such ‘sampling’ isn’t occurring (and so can’t be interfering with our assessment). (Ex: Consider the strange case where some bat enjoys a functionally isolated sonar experience.)

CONSTRAINTS ON RATIONALITY, NOT PHENOMENALITY

Objection: “What’s strange about these cases is not that they are *incoherent* but that they ask us to take up the point of view of a subject who suffers from an unfamiliar, severe form of irrationality. Strange subjects fail, in systematic ways, to draw the correct inferences from experience, or otherwise marshal the contents of their mind in service of cognitive tasks.”

Reply 1: This objection doesn’t avoid or account for the seeming difficulty of ascribing strange experiences to the subjects for whom they are inaccessible. So maybe collapses into bare denial...

Reply 2: You are rationally criticizable insofar as you fail to integrate all the information available to you in a systematic, consistent way. But in a strange experience case, the subject does *not* have access to key information, and so she has not even been given a chance respond to and act on this information in a rational way. (Ex: A detective who never gains access to destroyed evidence.)

CONSTRAINTS ON PHENOMENAL SUBJECTS, NOT PHENOMENALITY

Objection: “Strange experience cases may reveal a priori constraints on phenomenal subjecthood, but such constraints shouldn’t be confused for constraints on phenomenal experience itself.”

Strange experiences don’t show that the phenomenal, as a conceptual domain, must be understood in functional terms; rather, they just show that some particular functional arrangement of phenomenal experiences implicates functional facts. And it doesn’t directly challenge the idea that any arrangement of the atomic phenomenal facts is consistent with any set of functional facts.

Reply 1: Phenomenal-subject facts are (probably?) phenomenal facts. If “S is a P-conscious subject” a priori entails “S is A-conscious”, that still undermines Conceptual Dualism.

Whether some set of phenomenal experiences is unified together in, or belongs to some phenomenal subject should count as a phenomenal fact. That is, *that S is an experiencer* seems to fall clearly on the phenomenal side of any phenomenal/functional line. So even if “S is a subject of phenomenal experience” a priori entails “S is A-conscious”, I’ll have successfully undermined

Conceptual Dualism. And once we break the seal, we may rightly question the motivations for keeping these conceptual domains separate in the first place.

Reply 2: It is not clear that we can conceive of individual experiences that do not belong to a subject. And so a functional constraint on phenomenal subject-hood would, consequently, constrain the entire phenomenal domain. (The only way to test this would be consider cases not unlike those we've already been considering: cases of subject-less experiences.)

Not clear that this objection actually represents a substantive argumentative move.

CONSTRAINTS ON CAUSAL CONNECTION, NOT ACCESS

Objection: "Strange experience cases maybe show that a subject's phenomenal states can't be causally isolated from each other—but that's a very minimal, uninteresting constraint. And it falls short of an *access* constraint."

Reply 1: Even this minimal constraint violates Conceptual Dualism. (c.f. Zombie cases don't reveal a 'minimal' phenomenal constraint.)

Reply 2: The kind of causal connection matters here. Your A-conscious states are, in some sense, causally connected to: subconscious processes, the digestive processes in your gut, to all external activity. And yet *mere* causal connection doesn't seem to address the difficulty brought out by strange experience cases. The causal connection has to be more direct. And that more direct sense is just what 'access' is meant to capture.

BARE DENIAL

Objection: "You're right that the above objections can't be used to decisively remove the challenge. But the conceivability of Strange Experiences is *basic*. When intuitions diverge over such basic conceptual 'starting points', all that's left is a familiar sort of bare disagreement. You're *just wrong*."

Same way that those who insist non-black squares are inconceivable are *just wrong*.

Reply: I agree that we should embrace the *basicness* of this issue. But even a stalemate would represent progress here, since Conceptual Dualism (in my estimation) hasn't received enough scrutiny. And we can still make progress on basic disagreements by drawing out the downstream consequences of each position, and seeing how those consequences interact with our other commitments. If demonstrating the full range of such consequences is not ultimately decisive, that would then support an enlightened agnosticism—it would not permit stubborn partisanship.

Comments/questions: jacksonkernion@gmail.com